

# Class 7: Machine Learning


AUTHOR

Nichelle Camden

## K-means Clustering

Let's make up some data to cluster.

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))  
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



The function to do k-means clustering in base R is called `kmeans()`. We give this our input data for clustering and the number of clusters we want `centers`.

```
km <- kmeans(x, centers = 4, nstart = 20)  
km
```

K-means clustering with 4 clusters of sizes 17, 13, 17, 13

Cluster means:

	x	y
1	2.535453	-3.972393
2	-2.297487	3.892280
3	-3.972393	2.535453
4	3.892280	-2.297487

Clustering vector:

```
[1] 3 2 2 2 2 3 3 2 2 2 2 3 3 2 3 2 3 3 3 2 3 3 3 3 3 3 2 3 1 4 1 1 1 1 1 1
[39] 4 1 1 1 4 1 4 1 4 1 1 4 4 4 4 1 1 4 4 4 4 1
```

Within cluster sum of squares by cluster:

```
[1] 19.97669 14.04267 19.97669 14.04267
(between_SS / total_SS = 95.0 %)
```

Available components:

[1] "cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6] "betweenss"	"size"	"iter"	"ifault"	

```
km$cluster
```


```
[1] 3 2 2 2 2 3 3 2 2 2 2 3 3 2 3 2 3 2 3 3 3 2 3 3 3 3 3 2 3 1 4 1 1 1 1 1 1
[39] 4 1 1 1 4 1 4 1 4 1 1 4 4 4 4 1 1 4 4 4 4 1
```

```
km$centers
```

	x	y
1	2.535453	-3.972393
2	-2.297487	3.892280
3	-3.972393	2.535453
4	3.892280	-2.297487

Q. plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=1.5)
```



## Hierarchical Clusters

The `hclust()` function performs hierarchical clustering. The big advantage here is that I don't need to tell it "k" the number of clusters..

To run `hclust()` I need to provide a distance matrix as input (not the original data).

```
hc <- hclust( dist(x))  
hc
```


Call:

```
hclust(d = dist(x))
```

Cluster method : complete  
Distance : euclidean  
Number of objects: 60

```
plot(hc)  
abline(h=8, col="red", lty= 2)
```

## Cluster Dendrogram



To get my "main" result (cluster membership) I want to "cut" this tree to yield "branches" whos "leaves" are the members of the cluster.

```
cutree(hc, h=8)
```


```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

More often, we will use `cutree()` with `k=2` for example

```
grps <- cutree(hc, k=3)
```

Make a plot of our `hclust` results, i.e. our data colored by cluster assignment

```
plot(x, col=grps)
```



## Principal Component Analysis (PCA)

Read data for UK food trends from online

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494

Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

##Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
ncol(x)
```

```
[1] 4
```

```
nrow(x)
```

```
[1] 17
```

There are 5 columns and 17 rows. The food column counts as a column, in addition to the four countries. (before arguing with the url and changing the row names)

```
#View(x)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
x[,-1] #shows everything except the first row (food)
```

	Wales	Scotland	N.Ireland
Cheese	103	103	66
Carcass_meat	227	242	267
Other_meat	803	750	586
Fish	160	122	93
Fats_and_oils	235	184	209
Sugars	175	147	139
Fresh_potatoes	874	566	1033
Fresh_Veg	265	171	143
Other_Veg	570	418	355
Processed_potatoes	203	220	187
Processed_Veg	365	337	334
Fresh_fruit	1137	957	674
Cereals	1582	1462	1494
Beverages	73	53	47
Soft_drinks	1256	1572	1506

Alcoholic_drinks	475	458	135
Confectionery	64	62	41


```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267
685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139

**Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?**


I prefer to add a row.names argument when calling in the url so that the first row doesn't continually get deleted every time that chunk of code is run.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```






**Q3: Changing what optional argument in the above barplot() function results in the following plot?**

Beside = FALSE

**Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?**

If a point lies on the diagonal line I think that means it has the same value for both axis, so that value is the same or very similar for both countries represented in the plot. (the fold change is equal to 0)

```
pairs(x, col=rainbow(10), pch=16)
```



## Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Whatever category the dark blue dot represents seems to stand out for Northern Ireland in comparison to the other countries.

##PCA to the rescue!

The main function in base R to do PCA is called `prcomp()`. One issue with the `prcomp()` functions is that it expects the transpose of our data as input.


```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3
Standard deviation	379.8991	260.5533	1.515e-13
Proportion of Variance	0.6801	0.3199	0.000e+00
Cumulative Proportion	0.6801	1.0000	1.000e+00

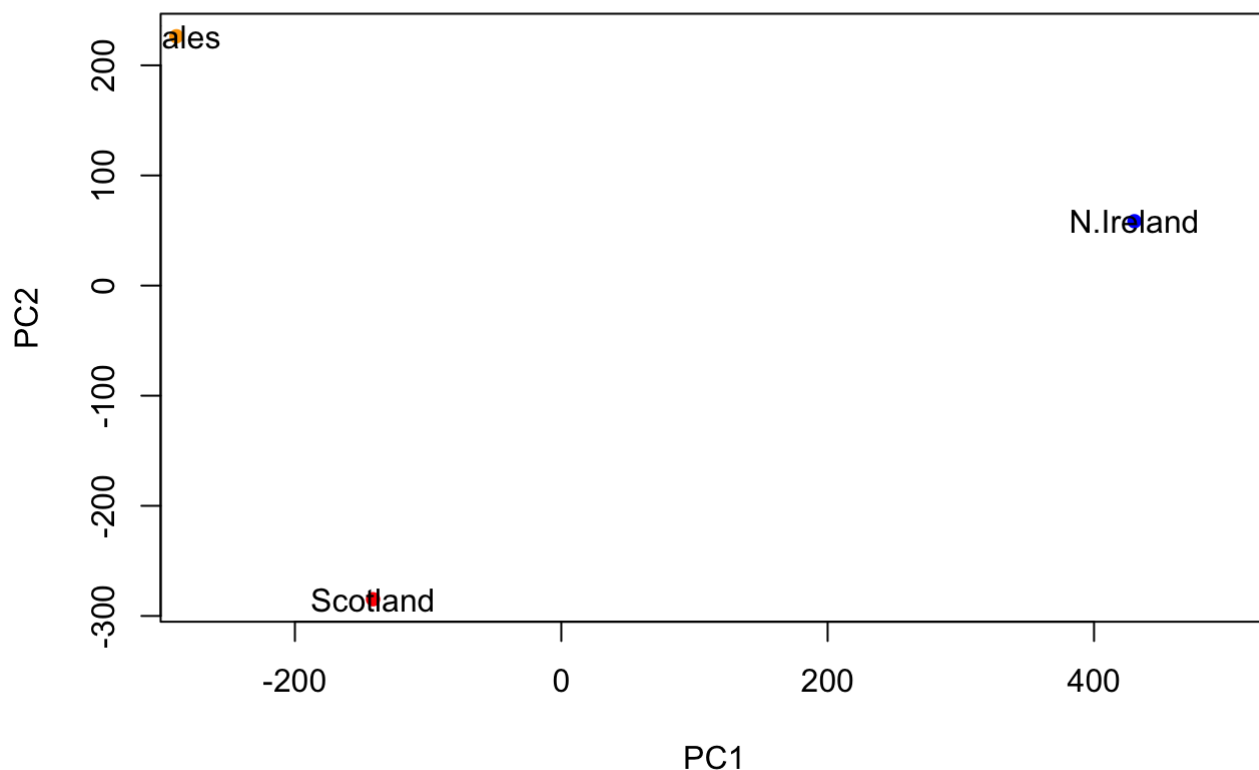
The object returned by `prcomp` has our results that include a `$x` component. This is our "scores" along the PCs (i.e. the plot of our data along the new PC axis)

```
plot(pca$x[,1], pca$x[,2],  
     xlab="PC1", ylab="PC2",  
     col=c("orange", "red", "blue", "darkgreen"),  
     pch=16)
```




**Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.**

```
# Plot PC1 vs PC2  
plot(pca$x[,1], pca$x[,2],  
     xlab="PC1", ylab="PC2",  
     xlim=c(-270,500), col=c("orange", "red", "blue", "darkgreen"), pch=16)  
text(pca$x[,1], pca$x[,2], colnames(x))
```




**Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.**

```
plot(pca$x[,1], pca$x[,2],  
     xlab="PC1", ylab="PC2",  
     xlim=c(-270,500),  
     col=c("orange", "red", "blue", "darkgreen"),)  
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "darkgreen"))
```



Lets focus on PC1 as it accounts for > 90% of variance

```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Northern Ireland consumes less fresh fruit, less alcoholic drinks, and more fresh potatoes than the other countries.