

Class06

AUTHOR
Nichelle Camden

PUBLISHED
October 14, 2022

Function basics

All functions in R have at least 3 things:

- A **name** (we pick this),
- Input **arguments** (there can be loads that are comma separated),
- A **body** (the R code that does the work).

```
# example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)

#and later we'll try it with this data
gradebook <- read.csv("student_homework.csv")
```

I can use the `mean()` function to get the average

```
mean(student1)
```

```
[1] 98.75
```

To find the lowest value, I can use the `min()` function

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

```
min(student1)
```

```
[1] 90
```

I found the `which.min()` function, what does it do?

```
which.min(student1)
```

```
[1] 8
```

The minimum value is in the 8th position for “student1”.

There are a few different ways to take that value out from the average calculation, but the index trick might be the easiest/most useful. (It could also work to sort the students' scores from lowest to highest, and then drop the first score for all.)

```
#lil refresh on the index trick:  
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

```
student1[-which.min(student1)]
```

```
[1] 100 100 100 100 100 100 100
```

Then I can take the mean

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

Will it work with "student2"?

```
student2[-which.min(student2)]
```

```
[1] 100 NA 90 90 90 90 97
```

Kind of, but not really. It just got rid of the lowest numerical value, but not the actual lowest value (the NA)

```
mean(student2[-which.min(student2)])
```

```
[1] NA
```

```
mean(student2, na.rm=T)
```

```
[1] 91
```

We need another way...

Can I replace NA values with zero? No homework submission = 0 try google

```
is.na(student2)
```

```
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE
```

```
student2[ is.na(student2) ] <- 0  
student2
```

```
[1] 100  0  90  90  90  90  97  80
```

```
c(T,T,F)
```

```
[1]  TRUE  TRUE FALSE
```

```
!c(T,T,F)
```

```
[1] FALSE FALSE  TRUE
```

```
# the ! flips it
```

```
is.na(student3)
```

```
[1] FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
student3[ is.na(student3) ] <- 0
student3
```

```
[1] 90  0  0  0  0  0  0  0
```

```
positions <- is.na(student2)
student2
```

```
[1] 100  0  90  90  90  90  97  80
```

```
student3[positions] <- 0
student3
```

```
[1] 90  0  0  0  0  0  0  0
```

```
#could also use "missing" instead of "positions"
```

```
student2[is.na(student2)] <- 0
mean(student2 [-which.min(student2)])
```

```
[1] 91
```

Re-write my snippet to be more simple for Q1

Q1

```
x <- student1
x[is.na(x)] <- 0
```

```
x[is.na(x)] <- 0  
mean(x[-which.min(x)])
```

[1] 100

```
#can replace "x" with student2 and 3  
x <- student2  
x[is.na(x)] <- 0  
mean(x[-which.min(x)])
```

[1] 91

```
x <- student3  
x[is.na(x)] <- 0  
mean(x[-which.min(x)])
```

[1] 12.85714

```
grade <- function(x) {  
  x[ is.na(x)] <- 0  
  mean(x[-which.min(x)])  
}
```

Now use that to grade student1 and others

```
grade(student1)
```

[1] 100

```
grade(student2)
```

[1] 91

```
grade(student3)
```

[1] 12.85714

Another way to do above is to highlight snippet -> Code (at top) -> Extract Function -> name the function (here I named it "grade")

```
grade <- function(x) {  
  x[is.na(x)] <- 0  
  mean(x[-which.min(x)])  
}
```

Who is the top scoring student overall in the gradebook?

(can get the data this way or the way I used above with the "read.csv("file_name") if it's downloaded already)

```
gradebook <- read.csv("https://tinyurl.com/gradeinput",
row.name = 1)
head(gradebook)
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77

Now I want to introduce the `apply()` function.

```
results <- apply(gradebook, 1, grade)
results
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

I can use `which.max` to find where the largest/ max value is in this results vector

```
which.max(results)
```

```
student-18
18
```

Q3

From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall)?

We can use `apply()` again, but this time over columns (use 2 instead of 1 so margin=2)

```
apply(gradebook, 2, sum, na.rm= TRUE)
```

hw1	hw2	hw3	hw4	hw5
1780	1456	1616	1703	1585

```
lowest_score <- apply(gradebook, 2, sum, na.rm= TRUE)
lowest_score
```

```
hw1 hw2 hw3 hw4 hw5
1780 1456 1616 1703 1585
```

I can use my eyeballs to see that homework 2 was the toughest, but I can also get R to tell me explicitly (incase datasets are too big in the future)

```
which.min(lowest_score)
```

```
hw2
2
```

Q4

Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)?

```
#cor(gradebook$hw1, results)
#cor(gradebook$hw2, results)

mask <- gradebook
mask[ is.na(mask)] <- 0
mask
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	0	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	0	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	0
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	0	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

```
cor(mask$hw5, results)
```

```
[1] 0.6325982
```

It looks like homework 5 is highly correlated, but let's use the `apply()` function over the masked gradebook so we don't have to retype hw1, hw2, etc

```
apply(mask, 2, cor, y=results)
```

	hw1	hw2	hw3	hw4	hw5
	0.4250204	0.1767780	0.3042561	0.3810884	0.6325982